

Predictive Modeling for Liver Disease Detection

Step 1: Problem Definition

This project develops a simple tool using common health data to predict liver disease early, helping doctors provide faster and better care. Liver disease often goes unnoticed until it becomes serious, as early symptoms are mild or absent. Diagnosing it quickly is challenging, especially where medical resources are limited. This project aims to develop a simple program that uses common health data—like age, weight, alcohol use, family history, and blood tests—to predict the risk of liver disease. The tool will help doctors detect liver problems early, enabling faster treatment, saving time and resources, and improving patient care.

This project aims to develop a machine learning model that predicts the presence of liver disease in patients based on various medical attributes such as age, gender, BMI, alcohol consumption, smoking status, genetic risk, physical activity, diabetes, hypertension, and liver function test results. Utilizing the <https://www.kaggle.com/datasets/rabieelkharoua/predict-liver-disease-1700-records-dataset/data> Dataset, we will train, evaluate, and interpret models to assist in clinical decision-making.

Why This Is Important:

- **Detect Early:** Find signs of liver disease before symptoms appear.
- **Save Lives:** Early treatment can stop the disease from getting worse.
- **Save Money:** Reduce expensive tests and hospital visits.
- **Help Doctors:** Provide support to doctors, especially when they have limited time or resources.
- **Reach Everyone:** Make liver disease detection easier in places with fewer medical facilities.

How This Can Be Used:

- Quickly identify patients who need more tests or care.
- Support doctors with clear, easy-to-understand risk predictions.
- Help prioritize patients based on how likely they are to have liver disease.
- Show which lifestyle habits or test results matter most for liver health.
- Be used in clinics, hospitals, and remote healthcare centers.

This project aims to make liver disease detection easier, faster, and more reliable, so more people get help early and stay healthier.

Adding Necessary Libraries For The Project

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
import xgboost as xgb
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import precision_score, recall_score, f1_score
import warnings
warnings.filterwarnings('ignore')
```

Step 2: Data Understanding

Loading and Analyzing the Data


Understanding the Dataset Columns

| Column Name | What It Means |
|--------------------|--|
| Age | How old the person is (20 to 80 years) |
| Gender | Male (0) or Female (1) |
| BMI | Body weight and height ratio (15 to 40) |
| AlcoholConsumption | How much alcohol the person drinks each week |
| Smoking | Does the person smoke? No (0) or Yes (1) |
| GeneticRisk | Family risk for liver disease: Low (0), Medium (1), High (2) |
| PhysicalActivity | Hours of exercise per week (0 to 10) |
| Diabetes | Does the person have diabetes? No (0), Yes (1) |
| Hypertension | Does the person have high blood pressure? No (0), Yes (1) |
| LiverFunctionTest | Liver health test score (20 to 100) |
| Diagnosis | Liver disease: No (0), Yes (1) |

```
In [2]: df = pd.read_csv('Liver_disease_data.csv')
df.head()
```

Out[2]:

| | Age | Gender | BMI | AlcoholConsumption | Smoking | GeneticRisk | PhysicalActivit |
|---|-----|--------|-----------|--------------------|---------|-------------|-----------------|
| 0 | 58 | 0 | 35.857584 | 17.272828 | 0 | 1 | 0.65894 |
| 1 | 71 | 1 | 30.732470 | 2.201266 | 0 | 1 | 1.67055 |
| 2 | 48 | 0 | 19.971407 | 18.500944 | 0 | 0 | 9.92830 |
| 3 | 34 | 1 | 16.615417 | 12.632870 | 0 | 0 | 5.63012 |
| 4 | 62 | 1 | 16.065830 | 1.087815 | 0 | 1 | 3.56621 |



Dataset Shape

To check the number of rows and columns in the dataset, use the `.shape` property:

In [3]: `df.shape`

Out[3]: (1700, 11)

Checking for Missing Values

In [4]: `df.isnull().sum()`

Out[4]:

| | |
|--------------------|---|
| Age | 0 |
| Gender | 0 |
| BMI | 0 |
| AlcoholConsumption | 0 |
| Smoking | 0 |
| GeneticRisk | 0 |
| PhysicalActivity | 0 |
| Diabetes | 0 |
| Hypertension | 0 |
| LiverFunctionTest | 0 |
| Diagnosis | 0 |

dtype: int64

Checking for Duplicate Values

In [5]: `df.duplicated().sum()`

Out[5]: np.int64(0)

Dataset Summary:

Use the following command to generate a statistical summary of all **numerical features**:

In [6]: `df.describe()`

Out[6]:

| | Age | Gender | BMI | AlcoholConsumption | Smoking | Genet |
|--------------|-------------|-------------|-------------|--------------------|-------------|--------|
| count | 1700.000000 | 1700.000000 | 1700.000000 | 1700.000000 | 1700.000000 | 1700.0 |
| mean | 50.394118 | 0.504118 | 27.699801 | 9.832309 | 0.291765 | 0.5 |
| std | 17.641915 | 0.500130 | 7.210400 | 5.757472 | 0.454708 | 0.6 |
| min | 20.000000 | 0.000000 | 15.004710 | 0.003731 | 0.000000 | 0.0 |
| 25% | 35.000000 | 0.000000 | 21.455414 | 4.841811 | 0.000000 | 0.0 |
| 50% | 51.000000 | 1.000000 | 27.925367 | 9.828195 | 0.000000 | 0.0 |
| 75% | 66.000000 | 1.000000 | 33.957668 | 14.871671 | 1.000000 | 1.0 |
| max | 80.000000 | 1.000000 | 39.992845 | 19.952456 | 1.000000 | 2.0 |

Quick Data Overview:

Use the follwing command to see a quick summary of your dataset

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1700 entries, 0 to 1699
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   1700 non-null   int64
1   Gender                1700 non-null   int64
2   BMI                   1700 non-null   float64
3   AlcoholConsumption    1700 non-null   float64
4   Smoking               1700 non-null   int64
5   GeneticRisk           1700 non-null   int64
6   PhysicalActivity       1700 non-null   float64
7   Diabetes              1700 non-null   int64
8   Hypertension          1700 non-null   int64
9   LiverFunctionTest     1700 non-null   float64
10  Diagnosis             1700 non-null   int64
dtypes: float64(4), int64(7)
memory usage: 146.2 KB
```

Dataset Features

The dataset includes numeric and categorical features to predict liver disease.

In [8]: `continuous_features = ['Age', 'BMI', 'AlcoholConsumption', 'PhysicalActivity', 'categorical_features = ['Gender', 'Smoking', 'GeneticRisk', 'Diabetes', 'Hyperte', target = ['Diagnosis']`

In [9]: `df[categorical_features].nunique()`

```
Out[9]: Gender          2
        Smoking         2
        GeneticRisk     3
        Diabetes        2
        Hypertension    2
        dtype: int64
```

Data Understanding Observations

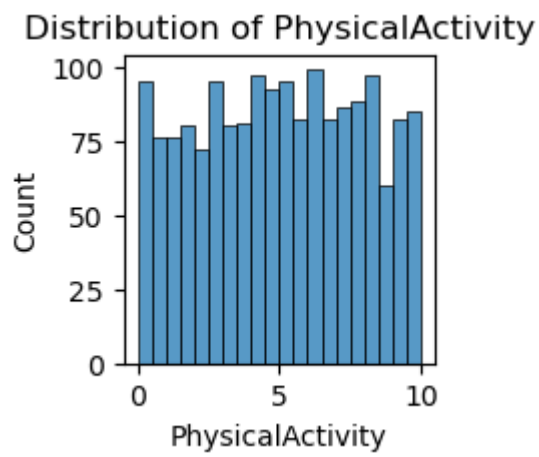

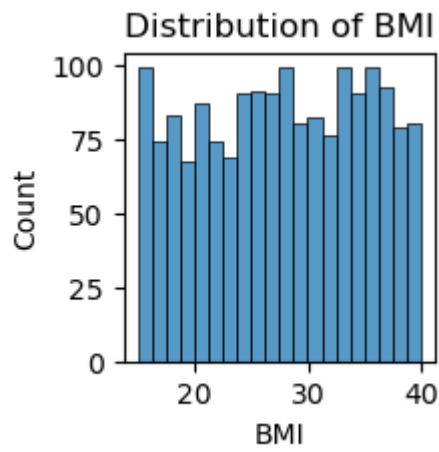

- **Dataset Size:** Contains 1700 records and 11 features, which is sufficient for analysis and modeling.
- **No Missing Values:** All features, including the target, have complete data with no missing values, so no data imputation is necessary.
- **No Duplicate Records:** The dataset does not contain any duplicate rows, ensuring data quality and integrity.
- **Feature Types:**
 - Continuous features: Age, BMI, AlcoholConsumption, PhysicalActivity, LiverFunctionTest
 - Categorical features: Gender, Smoking, GeneticRisk, Diabetes, Hypertension
- **Categorical Features:** Most categorical features are binary with 2 unique categories, except GeneticRisk which has 3 categories (low, medium, high).
- **Target Variable:** Diagnosis is binary, representing presence (1) or absence (0) of liver disease.
- **Overall Data Quality:** The dataset is clean, well-structured, and ready for the next steps of preprocessing and model building.

Step 3: Exploratory Data Analysis (EDA)

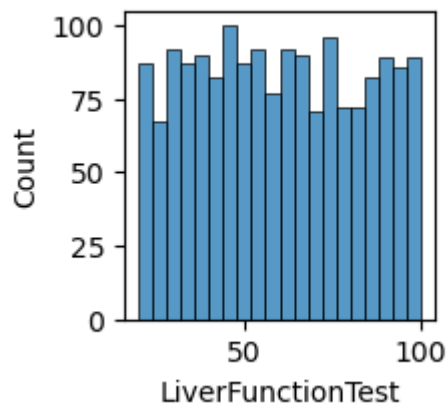
Visualize Distributions of Continuous Variables

```
In [10]: import seaborn as sns
import matplotlib.pyplot as plt

for feature in continuous_features:
    plt.figure(figsize=(2, 2))
    sns.histplot(df[feature], bins=20, stat="count")
    plt.title(f'Distribution of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.show()
```



Distribution of LiverFunctionTest




Visualize Categorical Variables


```
In [11]: import seaborn as sns
import matplotlib.pyplot as plt

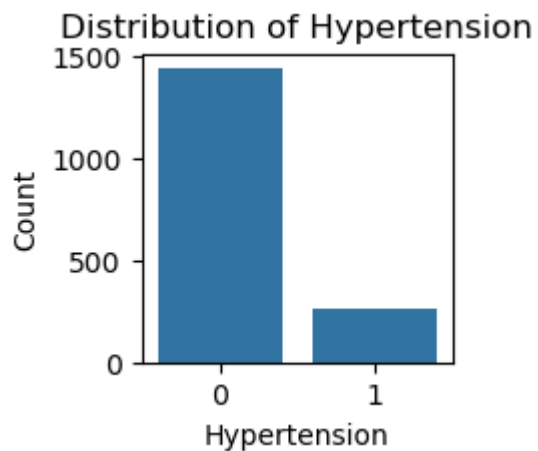

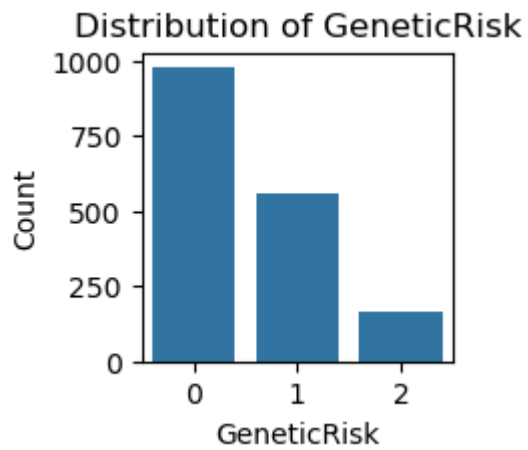
for feature in categorical_features:
    plt.figure(figsize=(2, 2))
    sns.countplot(x=df[feature])
    plt.title(f'Distribution of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.show()
```

Distribution of Gender



Distribution of Smoking





```
In [12]: from IPython.display import display, Markdown
for x in categorical_features:
    display(Markdown(f"##### Total count for {x} #####"))
    print(df[x].value_counts())
```

Total count for Gender

Gender

1 857

0 843

Name: count, dtype: int64

Total count for Smoking

Smoking

0 1204

1 496

Name: count, dtype: int64

Total count for GeneticRisk


```
GeneticRisk
0    978
1    557
2    165
Name: count, dtype: int64
```

Total count for Diabetes


```
Diabetes
0    1458
1     242
Name: count, dtype: int64
```

Total count for Hypertension

```
Hypertension
0    1437
1     263
Name: count, dtype: int64
```

Analyze Target Variable

```
In [13]: sns.countplot(x='Diagnosis', data=df, stat='count')
plt.title('Proportion of Liver Disease Diagnosis')
plt.ylabel('Probability')
plt.show()
```




```
In [14]: df["Diagnosis"].value_counts()
```

```
Out[14]: Diagnosis
1     936
0     764
Name: count, dtype: int64
```

Correlation Analysis

Correlation heatmap for continuous variables. Check if any features are highly correlated (multicollinearity).

```
In [15]: plt.title('Correlation Matrix (Continuous Variables)')
sns.heatmap(df[continuous_features].corr(), annot=True, cmap='coolwarm')
plt.show()
```



Observations: Correlation Between Continuous Variables

We checked correlation between continuous features only.

Important Correlation Values:

- **BMI vs LiverFunctionTest** → **+0.044** (highest, but still very weak)
- **AlcoholConsumption vs LiverFunctionTest** → **-0.043**
- **Age vs BMI** → **+0.038**
- All values are very small → no strong relation.
- So, we can use all features together no multicollinearity exists.

Step 4 : Data Wrangling

(4a)Checking For Skewness

Skewness shows whether the data is symmetrical, left-skewed, or right-skewed.

Conditions for Skewness

| Skewness Value | Interpretation |
|----------------|--------------------------------------|
| 0 | Perfectly symmetrical |
| > 0 | Right-skewed (tail on the right) |
| < 0 | Left-skewed (tail on the left) |
| -0.5 to +0.5 | Fairly symmetrical |
| -1 to -0.5 | Moderate left skew |
| +0.5 to +1 | Moderate right skew |
| < -1 or > +1 | Highly skewed (needs transformation) |

```
In [16]: df.skew()
```

```
Out[16]: Age                -0.040808
Gender                -0.016486
BMI                  -0.071939
AlcoholConsumption    0.018232
Smoking               0.916986
GeneticRisk           0.906531
PhysicalActivity      -0.023409
Diabetes              2.048946
Hypertension          1.911372
LiverFunctionTest     0.040151
Diagnosis             -0.203576
dtype: float64
```

Apply The BoxCox for Highly skewed column ie.Diabetes and Hypertension

```
In [17]: from scipy.stats import boxcox
cols_for_boxcox = ['Diabetes', 'Hypertension']
for col in cols_for_boxcox:
    if (df[col] <= 0).any():
        print(f"Skipping {col} – contains zero or negative values.")
        continue
    df[col], _ = boxcox(df[col])
    print(f"Applied Box-Cox on {col}.")
```

Skipping Diabetes – contains zero or negative values.

Skipping Hypertension – contains zero or negative values.

```
In [18]: df.skew()
```

```
Out[18]: Age          -0.040808
Gender        -0.016486
BMI           -0.071939
AlcoholConsumption  0.018232
Smoking        0.916986
GeneticRisk    0.906531
PhysicalActivity -0.023409
Diabetes       2.048946
Hypertension   1.911372
LiverFunctionTest 0.040151
Diagnosis      -0.203576
dtype: float64
```

(4b)Note: Here Encoding Is Not Required because all columns have numerical values

(4C).Scaling

- Scaling means changing big numbers into smaller ones so that all the features in the data are on a similar scale.
- This is important because machines sometimes treat bigger numbers as more important.
- If one feature has large values and another has small values, the model might focus more on the big ones.
- To avoid this, we scale the data so that all features are treated equally.
- Scaling should only be used on continuous data (like age, bmi,AlcoholConsumption and so on.)not on data with fixed categories or small whole numbers.
- If we apply scaling before splliting it will cause data leakage so apply during train test split

```
In [19]: df.columns
```

```
Out[19]: Index(['Age', 'Gender', 'BMI', 'AlcoholConsumption', 'Smoking', 'GeneticRisk',
               'PhysicalActivity', 'Diabetes', 'Hypertension', 'LiverFunctionTest',
               'Diagnosis'],
              dtype='object')
```

Step 5: Modelling

```
In [20]: X = df.drop(columns=['Diagnosis']) # Independent variables
y = df['Diagnosis'] # Target Variable
```

```
In [21]: X
```

Out[21]:

| | Age | Gender | BMI | AlcoholConsumption | Smoking | GeneticRisk | PhysicalAct |
|------|-----|--------|-----------|--------------------|---------|-------------|-------------|
| 0 | 58 | 0 | 35.857584 | 17.272828 | 0 | 1 | 0.65 |
| 1 | 71 | 1 | 30.732470 | 2.201266 | 0 | 1 | 1.67 |
| 2 | 48 | 0 | 19.971407 | 18.500944 | 0 | 0 | 9.92 |
| 3 | 34 | 1 | 16.615417 | 12.632870 | 0 | 0 | 5.63 |
| 4 | 62 | 1 | 16.065830 | 1.087815 | 0 | 1 | 3.56 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1695 | 42 | 0 | 38.498295 | 14.384688 | 0 | 1 | 0.99 |
| 1696 | 40 | 0 | 27.600094 | 5.431009 | 0 | 0 | 8.39 |
| 1697 | 38 | 0 | 38.730017 | 6.324302 | 1 | 2 | 9.31 |
| 1698 | 67 | 0 | 35.820798 | 16.899417 | 0 | 2 | 3.22 |
| 1699 | 80 | 0 | 24.060783 | 9.526447 | 0 | 0 | 9.26 |

1700 rows × 10 columns



In [22]:

y

Out[22]:

```
0      1
1      1
2      0
3      1
4      1
..
1695   1
1696   1
1697   1
1698   1
1699   1
```

Name: Diagnosis, Length: 1700, dtype: int64

We are splitting the dataset into:

- **80% training data**
- **20% testing data**

In [23]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.8,random_state=
```

In [24]:

```
columns_to_scale = ['Age', 'BMI', 'AlcoholConsumption', 'GeneticRisk', 'PhysicalAct']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Create a copy so original data is safe
X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()
# Initialize the scaler
scaler = StandardScaler() # important for joblib (varname)
# Fit on training data and transform both train and test
```

```
X_train_scaled[columns_to_scale] = scaler.fit_transform(X_train[columns_to_scale])
X_test_scaled[columns_to_scale] = scaler.transform(X_test[columns_to_scale])
```

```
In [25]: #-----Creating a DataFrame that stores all the metrics and performance of each
algorithms = ['logistic_Model', 'knn_Model', 'svm_Model', 'dt_Model', 'rf_Model']
metrics = ['TrainAccuracy', 'TestAccuracy', 'TrainPrecision', 'TestPrecision', 'TrainRecall', 'TestRecall', 'TrainF1', 'TestF1', 'CV']

analysis_df = pd.DataFrame(index=algorithms, columns=metrics)
```

```
In [26]: #-----DataFrame to store metrics useful for further analysis and Model Selection
analysis_df
```

```
Out[26]:
```

| | TrainAccuracy | TestAccuracy | TrainPrecision | TestPrecision | TrainRecall | TestRecall |
|-----------------------|---------------|--------------|----------------|---------------|-------------|------------|
| logistic_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| knn_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| svm_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| dt_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| rf_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| ada_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| gb_Model | NaN | NaN | NaN | NaN | NaN | NaN |
| xg_Model | NaN | NaN | NaN | NaN | NaN | NaN |

```
In [27]: #---Function that calculates all the metrics and Classification report and updates the DataFrame
def model_performance(model_key, model_obj, X_train, y_train, X_test, y_test, analysis_df):
    y_train_pred = model_obj.predict(X_train)
    y_test_pred = model_obj.predict(X_test)
```

```
    analysis_df.loc[model_key, 'TrainAccuracy'] = accuracy_score(y_train, y_train_pred)
    analysis_df.loc[model_key, 'TestAccuracy'] = accuracy_score(y_test, y_test_pred)
    analysis_df.loc[model_key, 'TrainPrecision'] = precision_score(y_train, y_train_pred)
    analysis_df.loc[model_key, 'TestPrecision'] = precision_score(y_test, y_test_pred)
    analysis_df.loc[model_key, 'TrainRecall'] = recall_score(y_train, y_train_pred)
    analysis_df.loc[model_key, 'TestRecall'] = recall_score(y_test, y_test_pred)
    analysis_df.loc[model_key, 'TrainF1'] = f1_score(y_train, y_train_pred)
    analysis_df.loc[model_key, 'TestF1'] = f1_score(y_test, y_test_pred)
```

```
    cv_score = cross_val_score(model_obj, X_train, y_train, cv=5, scoring='accuracy')
    analysis_df.loc[model_key, 'CV'] = cv_score
```

```
    print(f'■ Classification Report - {model_key} (Train)')
    print(classification_report(y_train, y_train_pred))
    print(f'■ Classification Report - {model_key} (Test)')
    print(classification_report(y_test, y_test_pred))
```

```
    # Confusion Matrix - Train
```

```
    cm_train = confusion_matrix(y_train, y_train_pred)
    disp_train = ConfusionMatrixDisplay(confusion_matrix=cm_train)
    disp_train.plot(cmap='Reds')
    plt.title(f'{model_key} - Confusion Matrix (Train)')
```

```
plt.show()

# Confusion Matrix - Test
cm_test = confusion_matrix(y_test, y_test_pred)
disp_test = ConfusionMatrixDisplay(confusion_matrix=cm_test)
disp_test.plot(cmap='Greens')
plt.title(f'{model_key} - Confusion Matrix (Test)')
plt.show()

return analysis_df
```

LOGISTIC REGRESSION

Modelling

Logistic Regression(Base Line Model)>>>>>>>>>>>>>>>>> 1st model

```
from sklearn.linear_model import LogisticRegression
Lr = LogisticRegression()
Lr.fit(X_train_scaled,y_train)
```

▼ LogisticRegression ⓘ ?

```
LogisticRegression()
```

```
Lr.coef_
```

```
array([[ 0.53261997,  1.21574022,  0.59059964,  1.3628255 ,  1.69515276,
         0.42220957, -0.39203883,  0.93703769,  1.45160119,  1.24841574]])
```

```
Lr.intercept_
```

```
array([-1.00791307])
```

```
s1=pd.DataFrame(Lr.predict_proba(X_train_scaled))
s1.drop(columns=[0],inplace=True)
s1
```

```
Out[31]:
```

| | 1 |
|------|----------|
| 0 | 0.026075 |
| 1 | 0.591947 |
| 2 | 0.066010 |
| 3 | 0.067497 |
| 4 | 0.328862 |
| ... | ... |
| 1355 | 0.707397 |
| 1356 | 0.041002 |
| 1357 | 0.577648 |
| 1358 | 0.066515 |
| 1359 | 0.977866 |

1360 rows × 1 columns

```
In [32]: Lr.predict(X_train_scaled)
```

```
Out[32]: array([0, 1, 0, ..., 1, 0, 1])
```

```
In [33]: Lr.predict_proba(X_train_scaled)
```

```
Out[33]: array([[0.97392461, 0.02607539],
                [0.40805257, 0.59194743],
                [0.93399013, 0.06600987],
                ...,
                [0.42235185, 0.57764815],
                [0.93348453, 0.06651547],
                [0.02213448, 0.97786552]])
```

Evaluation


```
In [34]: ypred_train = Lr.predict((X_train_scaled))
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(Lr,X_train_scaled,y_train,ypred_test= Lr.predict(X_test_scaled))
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

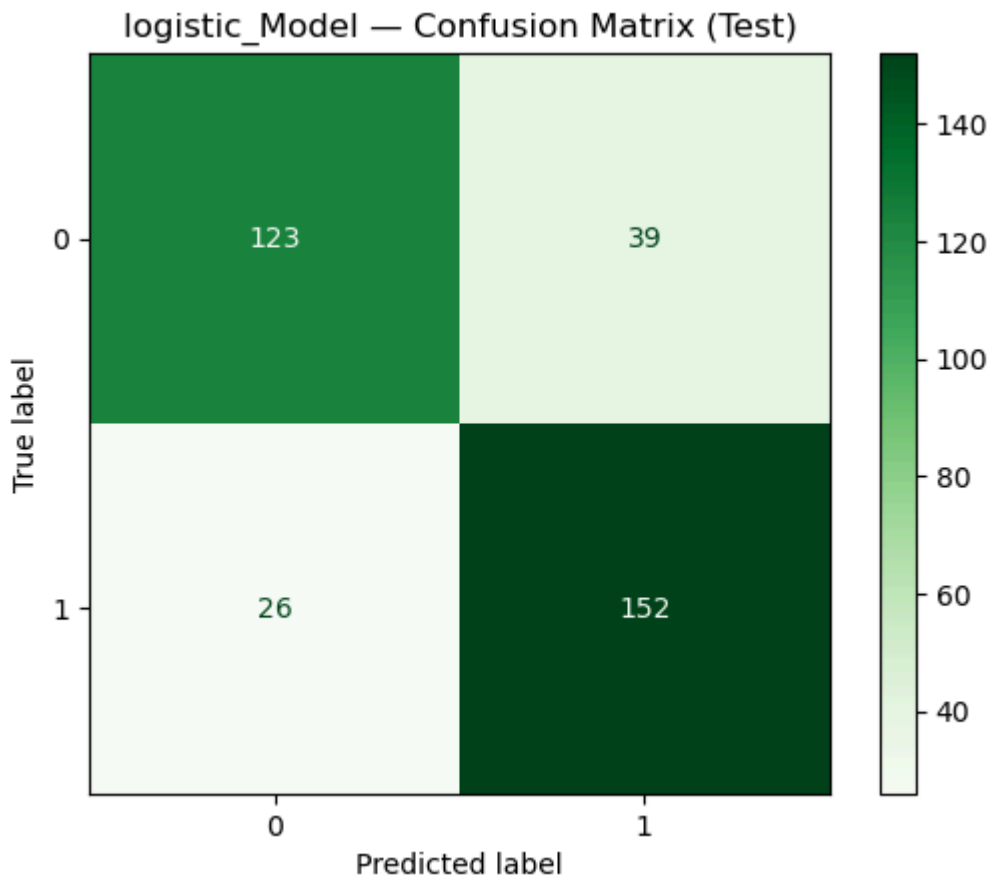
```
TRAIN ACCURACY 0.8323529411764706
THE CV SCORE(accuracy of model) 0.8286764705882353
TEST ACCURACY 0.8088235294117647
```

```
In [35]: logistic_Model_Report = model_performance('logistic_Model', Lr, X_train_scaled,
```


| Classification Report – logistic_Model (Train) | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.82 | 0.80 | 0.81 | 602 |
| 1 | 0.84 | 0.86 | 0.85 | 758 |
| accuracy | | | 0.83 | 1360 |
| macro avg | 0.83 | 0.83 | 0.83 | 1360 |
| weighted avg | 0.83 | 0.83 | 0.83 | 1360 |

| Classification Report – logistic_Model (Test) | | | | |
|---|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.83 | 0.76 | 0.79 | 162 |
| 1 | 0.80 | 0.85 | 0.82 | 178 |
| accuracy | | | 0.81 | 340 |
| macro avg | 0.81 | 0.81 | 0.81 | 340 |
| weighted avg | 0.81 | 0.81 | 0.81 | 340 |





KNN CLASSIFIER

Modelling

```
In [36]: from sklearn.neighbors import KNeighborsClassifier
estimator= KNeighborsClassifier()
param_grid={"n_neighbors": list(range(1,100))}
from sklearn.model_selection import GridSearchCV
cv_classifier=GridSearchCV(estimator,param_grid,cv=5,scoring='accuracy')
cv_classifier.fit(X_train_scaled,y_train)
cv_classifier.best_params_
```

Out[36]: {'n_neighbors': 9}

Evaluation

```
In [37]: from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=9)
knn.fit(X_train_scaled,y_train)
ypred_train = knn.predict(X_train_scaled)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(knn,X_train_scaled,y_train,cv=5))
ypred_test= knn.predict(X_test_scaled)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

TRAIN ACCURACY 0.8801470588235294

THE CV SCORE(accuracy of model) 0.8191176470588235

TEST ACCURACY 0.8058823529411765

```
In [38]: knn_Model_Report = model_performance('knn_Model', knn, X_train_scaled, y_train,
```

```
Classification Report - knn_Model (Train)
      precision    recall  f1-score   support


     0       0.84      0.90      0.87        602
     1       0.91      0.87      0.89        758


 accuracy          0.88          1360
 macro avg       0.88      0.88      0.88          1360
 weighted avg    0.88      0.88      0.88          1360
```

```
Classification Report - knn_Model (Test)
      precision    recall  f1-score   support

     0       0.78      0.82      0.80        162
     1       0.83      0.79      0.81        178

 accuracy          0.81          340
 macro avg       0.81      0.81      0.81          340
 weighted avg    0.81      0.81      0.81          340
```





SUPPORT VECTOR MACHINE

Modelling

FIRST TRY WITH DEFAULT PARAMS

```
In [39]: from sklearn.svm import SVC
svm = SVC(C=1, kernel="rbf")
svm.fit(X_train_scaled, y_train)
ypred_train = svm.predict(X_train_scaled)
print("TRAIN ACCURACY ", accuracy_score(y_train, ypred_train))
print("THE CV SCORE (accuracy of model)", cross_val_score(svm, X_train_scaled, y_train, cv=5, scoring='accuracy'))
ypred_test = svm.predict(X_test_scaled)
print("TEST ACCURACY", accuracy_score(y_test, ypred_test)) # Default value of C is 1
```

```
TRAIN ACCURACY 0.8852941176470588
THE CV SCORE (accuracy of model) 0.8522058823529411
TEST ACCURACY 0.8823529411764706
```

Hyperparameter Tuning For Svm Classifier

```
In [40]: from sklearn.model_selection import GridSearchCV
estimator = SVC()
param_grid = {"C": [0.001, 0.01, 0.1, 1, 10, 100], "kernel": ["linear", "rbf", "sigmoid"]}
grid = GridSearchCV(estimator, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train_scaled, y_train)
grid.best_params_
```

```
Out[40]: {'C': 1, 'kernel': 'rbf'}
```

Apply The SVM WITH BEST PARAMETERS

```
In [41]: svm = SVC(C=1, kernel="rbf")
svm.fit(X_train_scaled, y_train)
ypred_train = svm.predict(X_train_scaled)
print("TRAIN ACCURACY ", accuracy_score(y_train, ypred_train))
print("THE CV SCORE(accuracy of model)", cross_val_score(svm, X_train_scaled, y_train, cv=5))
ypred_test = svm.predict(X_test_scaled)
print("TEST ACCURACY", accuracy_score(y_test, ypred_test))
```

```
TRAIN ACCURACY 0.8852941176470588
THE CV SCORE(accuracy of model) 0.8522058823529411
TEST ACCURACY 0.8823529411764706
```

```
In [42]: svm_Model_Report = model_performance('svm_Model', svm, X_train_scaled, y_train,
```

```
■ Classification Report – svm_Model (Train)
      precision    recall  f1-score   support



    0       0.88      0.86      0.87         602
    1       0.89      0.91      0.90         758

 accuracy                   0.89         1360
 macro avg              0.88      0.88      0.88         1360
weighted avg              0.89      0.89      0.89         1360
```

```
■ Classification Report – svm_Model (Test)
      precision    recall  f1-score   support

    0       0.92      0.83      0.87         162
    1       0.86      0.93      0.89         178

 accuracy                   0.88         340
 macro avg              0.89      0.88      0.88         340
weighted avg              0.89      0.88      0.88         340
```



DECISION TREE

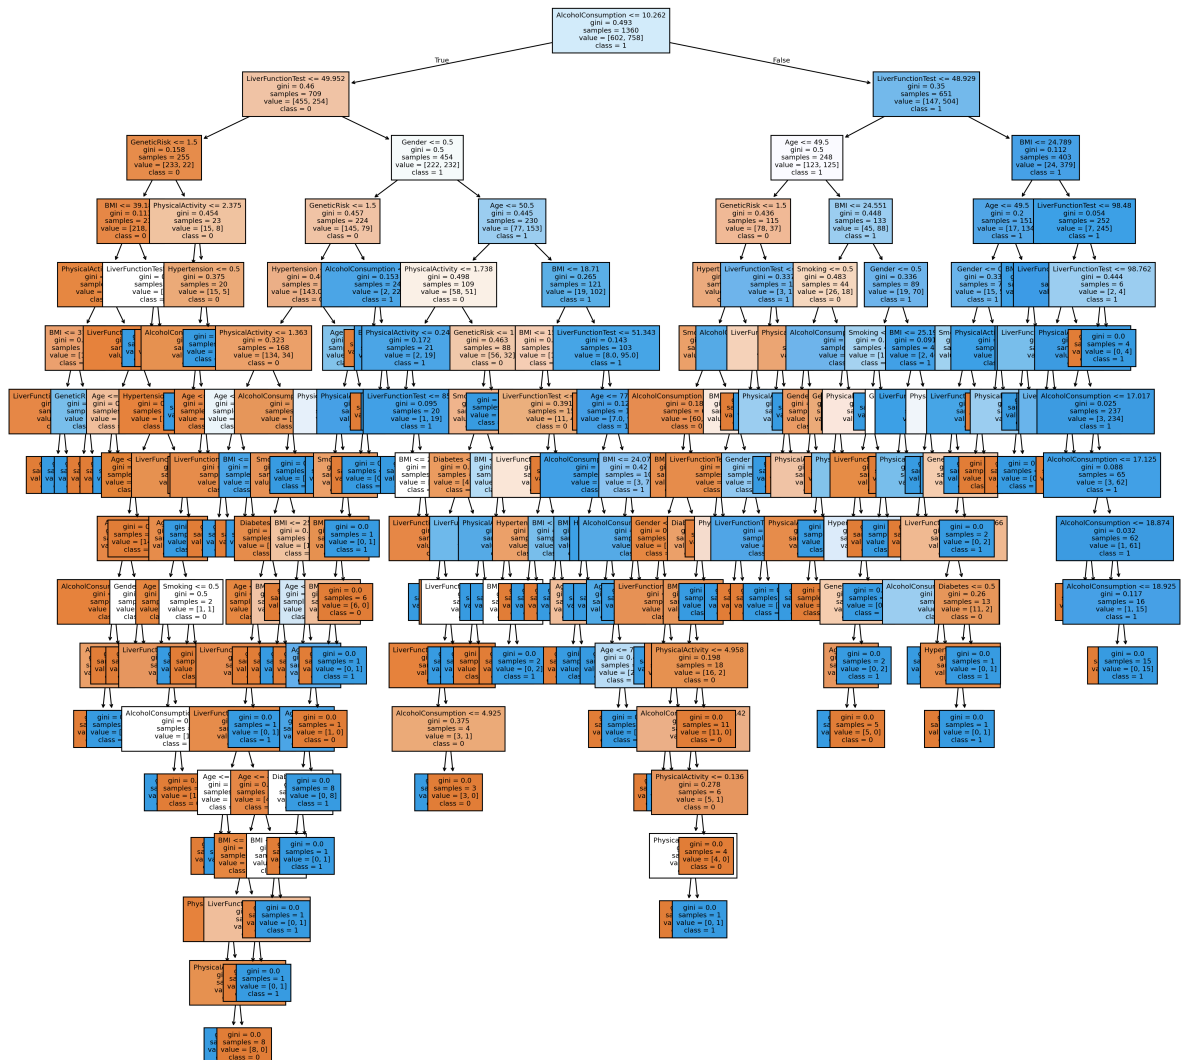
Modelling

FIRST TRY WITH DEFAULT PARAMS

```
In [43]: from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(random_state=42)
dt.fit(X_train,y_train) #while splliting it will calculate the gini
```

```
Out[43]: DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

```
In [44]: from sklearn.tree import plot_tree
plt.figure(figsize=(20,20),dpi=300)
plot_tree(dt,filled=True,feature_names=['Age', 'Gender', 'BMI', 'AlcoholConsumpt',
'PhysicalActivity', 'Diabetes', 'Hypertension', 'LiverFunctionTest'],clas
plt.show()
```



```
In [45]: ypred_train = dt.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(dt,X_train,y_train,cv=5
ypred_test= dt.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY 1.0
THE CV SCORE(accuracy of model) 0.8279411764705884
TEST ACCURACY 0.8382352941176471
```

In this data set there is Overfitting problem, then cut the tree using pruning which was given below

```
In [46]: estimator=DecisionTreeClassifier(random_state=42)
#params(which u want to tune and identify the best)
param_grid={"criterion":["gini","entropy"],"max_depth":[1,2,3,4,5,6]}
grid=GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
grid.fit(X_train,y_train)
grid.best_params_
```

```
Out[46]: {'criterion': 'gini', 'max_depth': 6}
```

```
In [47]: #best model
grid.best_estimator_
```

```
Out[47]: DecisionTreeClassifier
DecisionTreeClassifier(max_depth=6, random_state=42)
```

After creating decision tree model, using decision tree we can identify the important features

```
In [48]: grid.best_estimator_.feature_importances_
```

```
Out[48]: array([0.08041216, 0.06525154, 0.07028968, 0.2689609 , 0.04288555,
                0.07475267, 0.05941657, 0.          , 0.05394475, 0.28408618])
```

```
In [49]: s1=pd.DataFrame(index=X.columns,data=dt.feature_importances_,columns=["Feature Importance"])
```

```
Out[49]:
```

| | Feature Importance |
|--|--------------------|
|--|--------------------|

| | |
|--------------------|----------|
| Age | 0.102726 |
| Gender | 0.054350 |
| BMI | 0.110505 |
| AlcoholConsumption | 0.216515 |
| Smoking | 0.046592 |
| GeneticRisk | 0.060840 |
| PhysicalActivity | 0.096221 |
| Diabetes | 0.020850 |
| Hypertension | 0.046134 |
| LiverFunctionTest | 0.245267 |

```
In [50]: # Identify the important features
imp_columns=s1[s1["Feature Importance"] > 0].index.tolist()
imp_columns
```



```
Out[50]: ['Age',
          'Gender',
          'BMI',
          'AlcoholConsumption',
          'Smoking',
          'GeneticRisk',
          'PhysicalActivity',
          'Diabetes',
          'Hypertension',
          'LiverFunctionTest']
```

FINAL DECISION TREE MODEL

with best params and important columns

```
In [51]: X_imp=X[imp_columns]
X_train,X_test,y_train,y_test = train_test_split(X_imp,y,train_size=0.8,random_s
fdt=DecisionTreeClassifier(criterion='gini',max_depth=6,random_state=16)
fdt.fit(X_train,y_train)
ypred_train = fdt.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(fdt,X_train,y_train,cv=
ypred_test= fdt.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY 0.8911764705882353
THE CV SCORE(accuracy of model) 0.8301470588235293
TEST ACCURACY 0.8147058823529412
```

```
In [52]: dt_Model_Report = model_performance('dt_Model', fdt, X_train, y_train, X_test, y
```

```
■ Classification Report – dt_Model (Train)
      precision    recall  f1-score   support



    0       0.85       0.92       0.88         602
    1       0.93       0.87       0.90         758
```

```
   accuracy          0.89         1360
  macro avg       0.89       0.89       0.89         1360
weighted avg       0.89       0.89       0.89         1360
```

```
■ Classification Report – dt_Model (Test)
      precision    recall  f1-score   support

    0       0.79       0.84       0.81         162
    1       0.84       0.79       0.82         178

   accuracy          0.81         340
  macro avg       0.82       0.82       0.81         340
weighted avg       0.82       0.81       0.81         340
```



RANDOM FOREST

Modelling

```
In [53]: from sklearn.ensemble import RandomForestClassifier
rf=RandomForestClassifier(random_state=42)
rf.fit(X_train,y_train)
```

```
Out[53]: ▼      RandomForestClassifier
RandomForestClassifier(random_state=42)
```

FIRST TRY WITH DEFAULT PARAMS

```
In [54]: ypred_train = rf.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(rf,X_train,y_train,cv=5)
ypred_test= rf.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY  1.0
THE CV SCORE(accuracy of model) 0.8830882352941177
TEST ACCURACY  0.8970588235294118
```

control overfitting in Random Forest by limiting tree growth:

```
In [55]: estimator=RandomForestClassifier(random_state=42)
param_grid={'n_estimators' : list(range(1,50))}
grid=GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
grid.fit(X_train,y_train)
grid.best_params_
```

```
Out[55]: {'n_estimators': 42}
```

```
In [56]: #best model
grid.best_estimator_
```

```
Out[56]: ▼      RandomForestClassifier
RandomForestClassifier(n_estimators=42, random_state=42)
```

After creating random forest model we can identify the important features

```
In [57]: grid.best_estimator_.feature_importances_
```

```
Out[57]: array([0.11103888, 0.04173825, 0.12193691, 0.2468388 , 0.04150948,
0.04347761, 0.10299076, 0.01812926, 0.036712 , 0.23562806])
```

```
In [58]: s2=pd.DataFrame(index=X.columns,data=rf.feature_importances_,columns=["Feature I
s2
```

Out[58]:

| Feature Importance | |
|--------------------|----------|
| Age | 0.110756 |
| Gender | 0.043063 |
| BMI | 0.117660 |
| AlcoholConsumption | 0.250244 |
| Smoking | 0.041516 |
| GeneticRisk | 0.045610 |
| PhysicalActivity | 0.103177 |
| Diabetes | 0.018408 |
| Hypertension | 0.033062 |
| LiverFunctionTest | 0.236504 |

```
In [59]: # Identify the important features
imp_columns=s2[s2["Feature Importance"] > 0].index.tolist()
imp_columns
```

```
Out[59]: ['Age',
'Gender',
'BMI',
'AlcoholConsumption',
'Smoking',
'GeneticRisk',
'PhysicalActivity',
'Diabetes',
'Hypertension',
'LiverFunctionTest']
```

FINAL RANDOM FOREST MODEL

with best params and important columns


```
In [60]: X_imp=X[imp_columns]
X_train,X_test,y_train,y_test = train_test_split(X_imp,y,train_size=0.8,random_s
frf=RandomForestClassifier(n_estimators=44,random_state=16)
frf.fit(X_train,y_train)
ypred_train = frf.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(frf,X_train,y_train,cv=
ypred_test= frf.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```


```
TRAIN ACCURACY  0.9992647058823529
THE CV SCORE(accuracy of model) 0.8801470588235294
TEST ACCURACY  0.8852941176470588
```

```
In [61]: rf_Model_Report = model_performance('rf_Model', frf, X_train, y_train, X_test, y
```

| Classification Report – rf_Model (Train) | | | | | |
|--|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 1.00 | 1.00 | 1.00 | 602 | |
| 1 | 1.00 | 1.00 | 1.00 | 758 | |
| accuracy | | | 1.00 | 1360 | |
| macro avg | 1.00 | 1.00 | 1.00 | 1360 | |
| weighted avg | 1.00 | 1.00 | 1.00 | 1360 | |

| Classification Report – rf_Model (Test) | | | | | |
|---|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 0.91 | 0.85 | 0.88 | 162 | |
| 1 | 0.87 | 0.92 | 0.89 | 178 | |
| accuracy | | | 0.89 | 340 | |
| macro avg | 0.89 | 0.88 | 0.88 | 340 | |
| weighted avg | 0.89 | 0.89 | 0.88 | 340 | |





ADA BOOST

```
In [62]: from sklearn.ensemble import AdaBoostClassifier
ab=AdaBoostClassifier(random_state=42)
ab.fit(X_train,y_train)
```

```
Out[62]: ▼ AdaBoostClassifier ⓘ ?
AdaBoostClassifier(random_state=42)
```

FIRST TRY WITHOUT USING ANY PARAMS

```
In [63]: ypred_train = ab.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(ab,X_train,y_train,cv=5)
ypred_test= ab.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY 0.913235294117647
THE CV SCORE(accuracy of model) 0.8985294117647058
TEST ACCURACY 0.8823529411764706
```

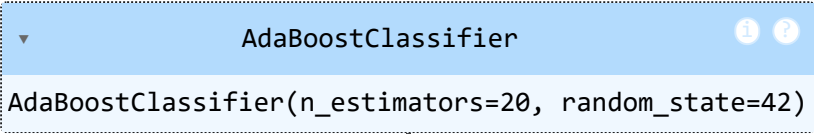
ADA BOOST>>> apply the HPT for identifying the best params

```
In [64]: estimator_ab=AdaBoostClassifier(random_state=42)
#params(which u want to tune and identify the best)
param_grid_ab={"n_estimators":list(range(1,51))} # click on shift tab(why 51 ta
grid=GridSearchCV(estimator_ab,param_grid_ab,scoring="accuracy",cv=5)
```

```
grid.fit(X_train,y_train)
grid.best_params_
```

Out[64]: {'n_estimators': 20}

In [65]: grid.best_estimator_

Out[65]: 

In [66]: grid.best_estimator_.feature_importances_

Out[66]: array([0.07196479, 0.06192163, 0.07273424, 0.19698819, 0.08619346,
0.08734884, 0.11936994, 0.05460205, 0.07289982, 0.17597704])

In [67]: s3=pd.DataFrame(index=X.columns,data=ab.feature_importances_,columns=["Feature I
s3

Out[67]:

| Feature Importance | |
|--------------------|----------|
| Age | 0.076920 |
| Gender | 0.043030 |
| BMI | 0.097278 |
| AlcoholConsumption | 0.169369 |
| Smoking | 0.069897 |
| GeneticRisk | 0.088627 |
| PhysicalActivity | 0.130709 |
| Diabetes | 0.057642 |
| Hypertension | 0.059329 |
| LiverFunctionTest | 0.207196 |

| Feature Importance | |
|--------------------|----------|
| Age | 0.076920 |
| Gender | 0.043030 |
| BMI | 0.097278 |
| AlcoholConsumption | 0.169369 |
| Smoking | 0.069897 |
| GeneticRisk | 0.088627 |
| PhysicalActivity | 0.130709 |
| Diabetes | 0.057642 |
| Hypertension | 0.059329 |
| LiverFunctionTest | 0.207196 |

In [68]: *# Identify the important features*
imp_columns=s3[s3["Feature Importance"] > 0].index.tolist()
imp_columns

Out[68]: ['Age',
'Gender',
'BMI',
'AlcoholConsumption',
'Smoking',
'GeneticRisk',
'PhysicalActivity',
'Diabetes',
'Hypertension',
'LiverFunctionTest']

Final Adaboost model with best hyperparameter and important columns

```
In [69]: X_imp=X[imp_columns]
X_train,X_test,y_train,y_test = train_test_split(X_imp,y,train_size=0.8,random_s
fab=AdaBoostClassifier(n_estimators=12)
fab.fit(X_train,y_train)
ypred_train = fab.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(fab,X_train,y_train,cv=
ypred_test= fab.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY 0.8948529411764706
THE CV SCORE(accuracy of model) 0.8948529411764706
TEST ACCURACY 0.8794117647058823
```

```
In [70]: ada_Model_Report = model_performance('ada_Model', fab, X_train, y_train, X_test,
```

```
■ Classification Report – ada_Model (Train)
      precision    recall  f1-score   support



      0       0.86       0.92       0.89         602
      1       0.93       0.88       0.90         758

 accuracy                   0.89         1360
 macro avg       0.89       0.90       0.89         1360
 weighted avg    0.90       0.89       0.90         1360

■ Classification Report – ada_Model (Test)
      precision    recall  f1-score   support

      0       0.89       0.85       0.87         162
      1       0.87       0.90       0.89         178

 accuracy                   0.88         340
 macro avg       0.88       0.88       0.88         340
 weighted avg    0.88       0.88       0.88         340
```

GRADIENT BOOST

```
In [71]: from sklearn.ensemble import GradientBoostingClassifier  
gb=GradientBoostingClassifier(random_state=42)
```

```
gb.fit(X_train,y_train)
```

```
Out[71]: ▾ GradientBoostingClassifier
GradientBoostingClassifier(random_state=42)
```

FIRST TRY WITHOUT USING ANY PARAMS

```
In [72]: ypred_train = gb.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(gb,X_train,y_train,cv=5)
ypred_test= gb.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY 0.9580882352941177
THE CV SCORE(accuracy of model) 0.8948529411764706
TEST ACCURACY 0.9088235294117647
```

GRADIENT BOOST>>> apply the HPT for identifying the best params

```
In [73]: estimator_gb=GradientBoostingClassifier(random_state=42)
#params(which u want to tune and identify the best)
param_grid_gb={"n_estimators":list(range(1,11))
               ,"learning_rate":[0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]}
grid=GridSearchCV(estimator_gb,param_grid_gb,scoring="accuracy",cv=5)
grid.fit(X_train,y_train)
grid.best_params_
```

```
Out[73]: {'learning_rate': 0.7, 'n_estimators': 10}
```

```
In [74]: grid.best_estimator_
```

```
Out[74]: ▾ GradientBoostingClassifier
GradientBoostingClassifier(learning_rate=0.7, n_estimators=10, random_state=42)
```

```
In [75]: grid.best_estimator_.feature_importances_
```

```
Out[75]: array([0.079026 , 0.06257073, 0.05550812, 0.30865613, 0.05807993,
               0.05049461, 0.05510406, 0.01631264, 0.03211812, 0.28212965])
```

```
In [76]: s4=pd.DataFrame(index=X.columns,data=gb.feature_importances_,columns=["Feature Importance"])
s4
```

Out[76]:

| Feature Importance | |
|--------------------|----------|
| Age | 0.080347 |
| Gender | 0.055993 |
| BMI | 0.071782 |
| AlcoholConsumption | 0.285440 |
| Smoking | 0.060968 |
| GeneticRisk | 0.057492 |
| PhysicalActivity | 0.062827 |
| Diabetes | 0.016049 |
| Hypertension | 0.033730 |
| LiverFunctionTest | 0.275373 |

```
In [77]: imp_columns=s4[s4["Feature Importance"] > 0].index.tolist()
imp_columns
```

```
Out[77]: ['Age',
'Gender',
'BMI',
'AlcoholConsumption',
'Smoking',
'GeneticRisk',
'PhysicalActivity',
'Diabetes',
'Hypertension',
'LiverFunctionTest']
```

Final Gradientboost model with best hyperparameter and important columns


```
In [78]: X_imp=X[imp_columns]
X_train,X_test,y_train,y_test = train_test_split(X_imp,y,train_size=0.8,random_s
fgb=GradientBoostingClassifier(n_estimators=10,learning_rate=0.7)
fgb.fit(X_train,y_train)
ypred_train = fgb.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(fgb,X_train,y_train,cv=
ypred_test= fgb.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```


```
TRAIN ACCURACY 0.9463235294117647
THE CV SCORE(accuracy of model) 0.8963235294117649
TEST ACCURACY 0.9205882352941176
```

```
In [79]: gb_Model_Report = model_performance('gb_Model', fgb, X_train, y_train, X_test, y
```

| Classification Report – gb_Model (Train) | | | | | |
|--|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 0.95 | 0.93 | 0.94 | 602 | |
| 1 | 0.94 | 0.96 | 0.95 | 758 | |
| accuracy | | | 0.95 | 1360 | |
| macro avg | 0.95 | 0.94 | 0.95 | 1360 | |
| weighted avg | 0.95 | 0.95 | 0.95 | 1360 | |

| Classification Report – gb_Model (Test) | | | | | |
|---|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 0.95 | 0.88 | 0.91 | 162 | |
| 1 | 0.90 | 0.96 | 0.93 | 178 | |
| accuracy | | | 0.92 | 340 | |
| macro avg | 0.92 | 0.92 | 0.92 | 340 | |
| weighted avg | 0.92 | 0.92 | 0.92 | 340 | |





EXTREME GRADIENTBOOSTING(XGBOOST)

```
In [80]: pip install xgboost
```

Requirement already satisfied: xgboost in c:\users\kaviti akhil\anaconda3\lib\site-packages (3.0.2)
Requirement already satisfied: numpy in c:\users\kaviti akhil\anaconda3\lib\site-packages (from xgboost) (2.1.3)
Requirement already satisfied: scipy in c:\users\kaviti akhil\anaconda3\lib\site-packages (from xgboost) (1.15.3)
Note: you may need to restart the kernel to use updated packages.

```
In [81]: from xgboost import XGBClassifier  
xgb = XGBClassifier(random_state=42)  
xgb.fit(X_train,y_train)
```

Out[81]:

```
XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               feature_weights=None, gamma=None, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=None, max_bin=None, max_cat_threshold=None,
               max_delta_step=None, max_leaf_nodes=None, min_child_weight=None,
               missing=None, monotone_constraints=None, multi_output_type='raw',
               num_parallel_tree=None, nthread=None, num_trees=None,
               print_eval_method=None, random_state=None,
               scale_pos_weight=None, seed=None, silent=None,
               subsample=None, tree_method=None, validate_features=None,
               verbosity=None, watchlog=None, weight_postfix=None,
               xgb_model=None, xgb_data_parallel_threshold=None,
               xgb_data_parallel_tree_method=None)
```

FIRST TRY WITHOUT USING ANY PARAMS

In [82]:

```
ypred_train = xgb.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(xgb,X_train,y_train,cv=5))
ypred_test= xgb.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))
```

```
TRAIN ACCURACY  1.0
THE CV SCORE(accuracy of model) 0.8875
TEST ACCURACY  0.8911764705882353
```

EXTREME GRADIENT BOOST>>> apply the HPT for identifying the best params

In [83]:

```
from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier
estimator_xgb = XGBClassifier()
param_grid_xgb = {
    "n_estimators": list(range(1, 11)),
    "learning_rate": [0, 0.1, 0.5, 1.0],
    "max_depth": [3, 4, 5],
    "gamma": [0, 0.15, 0.3, 0.5]
}
grid = GridSearchCV(estimator_xgb, param_grid_xgb, scoring="accuracy", cv=5)
grid.fit(X_train, y_train)
print("Best Parameters:", grid.best_params_)
```

```
Best Parameters: {'gamma': 0.5, 'learning_rate': 0.5, 'max_depth': 3, 'n_estimators': 10}
```

In [84]:

```
grid.best_estimator_.feature_importances_
```

Out[84]:

```
array([0.07172662, 0.11594246, 0.05710015, 0.17327702, 0.13394944,
       0.07379384, 0.07595403, 0.08205438, 0.07880992, 0.13739216],
      dtype=float32)
```

In [85]:

```
grid.best_estimator_
```

Out[85]:

```
XGBClassifier
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               feature_weights=None, gamma=0.5, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=0.5, max_bin=None, max_cat_threshold=None,
```

In [86]:

```
s4=pd.DataFrame(index=X.columns,data=xgb.feature_importances_,columns=["Feature Importance"])
```

Out[86]:

| Feature Importance | |
|--------------------|----------|
| Age | 0.067436 |
| Gender | 0.108262 |
| BMI | 0.049956 |
| AlcoholConsumption | 0.109963 |
| Smoking | 0.172660 |
| GeneticRisk | 0.117945 |
| PhysicalActivity | 0.049063 |
| Diabetes | 0.083409 |
| Hypertension | 0.145937 |
| LiverFunctionTest | 0.095369 |

In [87]:

```
imp_columns=s4[s4["Feature Importance"] > 0].index.tolist()
imp_columns
```

Out[87]:

```
['Age',
 'Gender',
 'BMI',
 'AlcoholConsumption',
 'Smoking',
 'GeneticRisk',
 'PhysicalActivity',
 'Diabetes',
 'Hypertension',
 'LiverFunctionTest']
```

Final ExtremeGradientboost model with best hyperparameter and important columns

In [88]:

```
X_imp=X[imp_columns]
X_train,X_test,y_train,y_test = train_test_split(X_imp,y,train_size=0.8,random_s
```

```

fxgb=XGBClassifier(n_estimators=10,learning_rate=0.5,max_depth=3,gamma=0.5)
fxgb.fit(X_train,y_train)
ypred_train = fxgb.predict(X_train)
print("TRAIN ACCURACY ",accuracy_score(y_train,ypred_train))
print("THE CV SCORE(accuracy of model)",cross_val_score(fxgb,X_train,y_train,cv
ypred_test= fxgb.predict(X_test)
print("TEST ACCURACY ",accuracy_score(y_test,ypred_test))

```

TRAIN ACCURACY 0.9345588235294118
 THE CV SCORE(accuracy of model) 0.9007352941176471
 TEST ACCURACY 0.8941176470588236

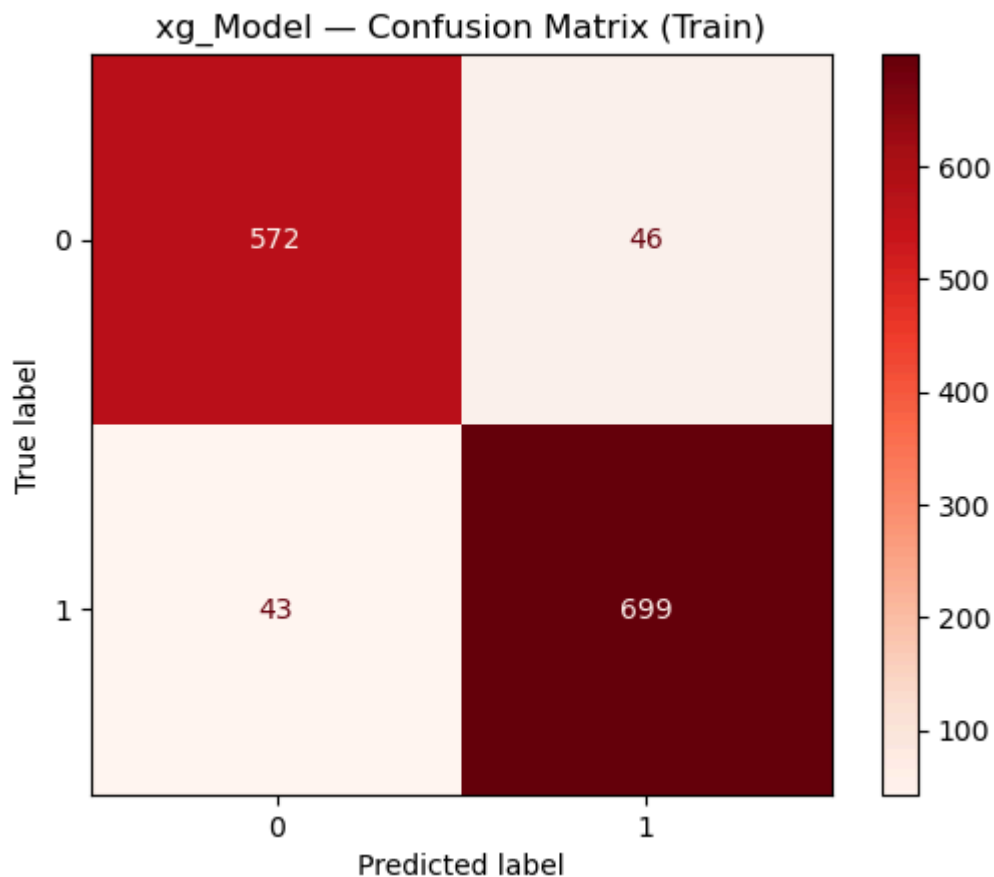
In [89]: `xg_Model_Report = model_performance('xg_Model', fxgb, X_train, y_train, X_test,`


■ Classification Report – xg_Model (Train)

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.93 | 0.93 | 0.93 | 618 |
| 1 | 0.94 | 0.94 | 0.94 | 742 |
| accuracy | | | 0.93 | 1360 |
| macro avg | 0.93 | 0.93 | 0.93 | 1360 |
| weighted avg | 0.93 | 0.93 | 0.93 | 1360 |

■ Classification Report – xg_Model (Test)

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.87 | 0.89 | 0.88 | 146 |
| 1 | 0.92 | 0.90 | 0.91 | 194 |
| accuracy | | | 0.89 | 340 |
| macro avg | 0.89 | 0.89 | 0.89 | 340 |
| weighted avg | 0.89 | 0.89 | 0.89 | 340 |





In [90]: analysis_df

Out[90]:

| | TrainAccuracy | TestAccuracy | TrainPrecision | TestPrecision | TrainRecall | T |
|----------------|---------------|--------------|----------------|---------------|-------------|---|
| logistic_Model | 0.832353 | 0.808824 | 0.841495 | 0.795812 | 0.861478 | |
| knn_Model | 0.880147 | 0.805882 | 0.913769 | 0.829412 | 0.866755 | |
| svm_Model | 0.885294 | 0.882353 | 0.890909 | 0.85567 | 0.905013 | |
| dt_Model | 0.891176 | 0.814706 | 0.929577 | 0.844311 | 0.870712 | |
| rf_Model | 0.999265 | 0.885294 | 1.0 | 0.867725 | 0.998681 | |
| ada_Model | 0.894853 | 0.879412 | 0.92887 | 0.87027 | 0.878628 | |
| gb_Model | 0.946324 | 0.920588 | 0.944228 | 0.895288 | 0.960422 | |
| xg_Model | 0.934559 | 0.894118 | 0.938255 | 0.915789 | 0.942049 | |

Here The best model is Gradient boost

In [91]: `from sklearn.ensemble import GradientBoostingClassifier`
`gb_Model=GradientBoostingClassifier(random_state=42)`
`gb_Model.fit(X_train,y_train)`

Out[91]: `GradientBoostingClassifier`
`GradientBoostingClassifier(random_state=42)`

Save The Model

```
In [92]: import joblib
# After training your model
joblib.dump(gb_Model, 'gboost_model.pkl')
```

```
Out[92]: ['gboost_model.pkl']
```

```
In [93]: import joblib
# After training your model
joblib.dump(scaler, 'scaler.pkl')
```

```
Out[93]: ['scaler.pkl']
```

Steps to be followed after saving the model

step1:Load the Model and Scaler

```
In [94]: import joblib
gb_Model = joblib.load('gboost_model.pkl') # your saved model
scaler = joblib.load('scaler.pkl') # your saved scaler
```

Step 2: Accept New Data (through user input)

```
In [95]: Age = int(input('Enter Age of the Patient (20 to 80): '))
Gender = int(input('Enter Gender (0 for Male, 1 for Female): '))
BMI = float(input('Enter BMI (e.g., 15 to 40): '))
AlcoholConsumption = float(input('Enter Alcohol Consumption per week (units): '))
Smoking = int(input('Does the patient smoke? (0 = No, 1 = Yes): '))
GeneticRisk = int(input('Enter Genetic Risk (0 = Low, 1 = Medium, 2 = High): '))
PhysicalActivity = float(input('Enter Hours of Physical Activity per week (0 to 10): '))
Diabetes = int(input('Does the patient have Diabetes? (0 = No, 1 = Yes): '))
Hypertension = int(input('Does the patient have Hypertension? (0 = No, 1 = Yes): '))
LiverFunctionTest = float(input('Enter Liver Function Test Score (20 to 100): '))

columns_to_scale = np.array([Age, BMI, AlcoholConsumption, GeneticRisk, PhysicalActivity, Diabetes, Hypertension, LiverFunctionTest])
scaled_columns = scaler.transform(columns_to_scale)
```

Step 3: Scale the New Data (only the scaled columns)

```
In [96]: #-----A 2D Array containing scaled values of the variables
print(scaled_columns)
# Age, BMI.....LiverFunctionTest
```

```
[[-1.6229822 -2.59317448 -1.35243455  0.72550143 -0.35591491 -0.64994161]]
```

```
In [97]: Age_scaled = scaled_columns[0][0]
        BMI_scaled = scaled_columns[0][1]
        AlcoholConsumption_scaled = scaled_columns[0][2]
        GeneticRisk_scaled = scaled_columns[0][3]
        PhysicalActivity_scaled = scaled_columns[0][4]
        LiverFunctionTest_scaled = scaled_columns[0][5]
```

```
In [98]: input_from_customer = [Age_scaled, Gender, BMI_scaled, AlcoholConsumption_scaled]
```

```
In [99]: input_array = np.array([input_from_customer])
```

Step 4: Make Predictions

```
In [100]: pred = gb_Model.predict(input_array)
         if pred[0] == 1:
             print("Liver Disease Detected")
         else:
             print("No Liver Disease")
```

No Liver Disease

Conclusion

We created a machine learning model to predict liver disease using patient data like age, BMI, alcohol consumption, and health indicators.

After testing 8 different models, **Gradient Boosting** gave the best performance with:

- **Test Accuracy:** 92.06%
- **Test Recall:** 96.07%
- **Test F1-Score:** 92.68%

It was the most balanced and reliable model, especially suitable for medical use cases where recall is important.

What We Did

- Selected important clinical and lifestyle features
 - Applied feature scaling on continuous variables
 - Trained 8 classification models and compared performance
 - Selected and saved the best model (Gradient Boosting)
 - Took user input to predict on new patient data
 - Displayed both prediction and probability
-

Use of This Project

This liver disease prediction system can help with:

- **Early detection** of liver disease in patients

- Use by doctors and hospitals as a **decision support tool**
 - Integration into **health apps** for quick risk checking
 - Use by individuals for **basic self-assessment**
-

This project shows how machine learning can support healthcare by providing fast, data-driven insights for better outcomes.